

Using a convolutional neural network to classify the Street View House Number dataset

Andrea Ferretti

andrea.ferretti1@studenti.unimi.it

The experiment described in this report consists of the development and use of a convolutional neural network to classify the Street View House Number dataset. In the first section a description of neural networks, particularly the kind of networks utilized in this experiment, is presented. In the second section the dataset and its peculiarities are introduced. In the third section is shown how the architecture of the model was developed and how the hyperparameters were tuned. In the fourth and final section the performance of the resulting model and some possible improvements are analyzed. The code used for this experiment can be found at <https://github.com/ferrettindr/cnn-svhn>.

Neural networks

Neural networks are a family of predictors characterized by the combination of simple computational units, called neurons. A neuron typically performs the following computation: $g(\mathbf{x}) = \sigma(\boldsymbol{\omega}^T \mathbf{x})$, where the elements of the vector $\boldsymbol{\omega}$ are the parameters of the neuron, σ is a non-linear function called activation function, and \mathbf{x} is the vector x_0, \dots, x_n where $x_0 = 1$ and x_i , for $i \in \{1, \dots, n\}$, is the output computed by another neuron. In the supervised learning setting, neurons are combined in a graph-like structure resulting in a computational network able to learn, by adjusting the parameters $\boldsymbol{\omega}$ of each neuron, the underlying mapping between data points and their corresponding labels.

One of the most basic architecture a neural network can assume is called feedforward network. In this type of network neurons are organized into successive layers: one input layer, one or more hidden layers, and one output layer. The computation flows from the input layer towards the output layer and each neuron passes its output to all the neurons in the next layer. The i -th neuron in the input layer simply outputs the value of the i -th dimension of the data

point. The neurons in the other layers compute a non-linear function of the weighted sum of the outputs of the neurons in the previous layer plus a bias; every neuron has therefore its own vector of weights and a bias term. The function computed by the neurons in the hidden layers is called the activation function and usually is some kind of sigmoid function such as the logistic function:

$$S(x) = \frac{1}{1 + e^{-x}} \quad (1)$$

The neurons in the output layer compute a function that varies depending on the type of problem: for a regressor the identity function can be used to obtain a real valued output for each neuron, for a classifier, instead, the softmax function is typically used. The softmax function $softmax : \mathbb{R}^m \rightarrow \mathbb{R}^m$ is defined as:

$$softmax(\mathbf{v})_i = \frac{e^{v_i}}{\sum_{t=0}^m e^{v_t}}, \text{ for } i = 0, \dots, m-1 \text{ and } \mathbf{v} \in \mathbb{R}^m \quad (2)$$

By having a network with as many output neurons as the number of possible classes to be predicted a probability distribution over those classes is obtained using softmax. It is to note that, in order for this function to be computed, every output neuron needs to know the values of the other output neurons (the various v_t) before the softmax is applied. Figure 1 gives an example of a feedforward neural network.

Consider the pair (\mathbf{x}, y) consisting of a data point and its corresponding label, let $\hat{\mathbf{y}}$ be the output of the network, in order to assess the goodness of the prediction a non-negative loss function $\ell(y, \hat{\mathbf{y}})$ is used so that the greater its value the worse is the prediction. In the case of a classification problem $y \in Y$, where Y is the ordered set of possible symbolic labels with $|Y| = m$, therefore, by having a network with m neurons in the output layer and using the *softmax* function, the network prediction $\hat{\mathbf{y}}$ represents a probability distribution over Y given the data point \mathbf{x} . For convenience of notation let y_i^* indicate the first element of Y for $i = 0$, the second element of Y for $i = 1$ and so on for all the elements of Y . This means that the i -th element of $\hat{\mathbf{y}}$ gives the probability that the label for the data point \mathbf{x} will be y_i^* in the network prediction. In this case, as the loss function, the categorical cross-entropy loss can be used and it is defined as:

$$\ell(y, \hat{\mathbf{y}}) = - \sum_{i=0}^{m-1} P(y = y_i^*) \log(\hat{y}_i) \quad (3)$$

where $P(y = y_i^*)$ is the probability that the true label for the data point \mathbf{x} is y_i^* . Since, in practice, to a single data point corresponds only one label with probability 1, equation (3) can




Figure 1: A feedforward network where on each neuron there is a label representing the value computed by it and passed to the nodes of the next layer through the arcs connecting them. A neuron assigns a weight ω to each incoming arch and performs the shown linear combination. The sigmoid function of equation (1) can be used as the activation function σ . The function computed by the neurons in the output layer could be *softmax* (2) where $\mathbf{v}_i = \omega_{y,i}^T \mathbf{n}_k + b_{y,i}$. Here the bias b is explicited in the computation rather than adding another dimension to the various vectors ω , \mathbf{n} , and \mathbf{x} so that $\omega_{r,i,0} = b_{r,i}$ for $r \in \{j, k, y\}$ and $x_0 = n_{p,0} = 1$ for $p \in \{j, k\}$.

be simplified to:

$$\ell(y, \hat{\mathbf{y}}) = - \sum_{i=0}^{m-1} \mathbb{I}(y = y_i^*) \log(\hat{y}_i) = -\log(\hat{y}_i), \text{ where } i \text{ verifies } y_i^* = y$$

where $\mathbb{I}(E) \in \{0, 1\}$ is the indicator function of an event E , meaning that $\mathbb{I}(E) = 1$ if and only if E occurs and $\mathbb{I}(E) = 0$ otherwise.

Given the loss function ℓ the network can be trained on the labeled data to minimize ℓ with respect to the parameters of each neuron. This is typically done with the mini-batched stochastic gradient descent algorithm, where, at each timestep, a fixed number of examples are randomly selected, creating a subset $B^{(t)}$ of the training set, and every parameter is iteratively updated following the rule:

$$\omega_{l,n,i}^{(t+1)} \leftarrow \omega_{l,n,i}^{(t)} - \eta^{(t)} \frac{1}{|B^{(t)}|} \sum_{b \in B^{(t)}} \frac{\partial \ell^{(t)}(y_b, \hat{\mathbf{y}}_b)}{\partial \omega_{l,n,i}} \quad (4)$$

where, as depicted in Figure 1, $\omega_{l,n,i}$ represents the i -th parameter of the n -th neuron in the

l -th layer, \hat{y}_b is the output computed by the network for the example (x_b, y_b) of the training set, and η_t is the learning rate that determines the update size. To calculate efficiently the partial derivatives with respect to all the parameters ω the backpropagation algorithm is used. The algorithm uses the chain rule to compute the derivative of a composite function:

$$\frac{df(g(x))}{dx} = \frac{df(g)}{dg} \frac{dg(x)}{dx}$$

When using a sigomid as the activation function of the neurons, the networks can experience the so called vanishing gradient problem, in which the gradients of the sigmoids of several neurons become so small that the series of multiplications of the chain rule causes, especially in deep networks, the update of the weight depending on those gradients to be almost non existant. This happens because the sigmoid function has small gradients for high absolute value inputs; one possible solutoin is to use the rectifier $f(x) = \max(0, x)$ as the activation function, this allows the gradient to always be equal to 1 when the output of the neuron is non-negative.

For input data that presents a grid-like structure a specific kind of network architecture called convolutional neural network (abbreviated ConvNet) has been developed. Similarly to feedforward networks it retains the forward direction of the computation, but it adds different types of layers such as convolutional layers and pooling layers. A convolutional layer is composed of several filters (also called kernels), each one performs a spatial convolution applying itself over the input dimensions; the output volume produced by this layer is called a feature map. A filter is composed by a mutlidimensional vector, where each element is a real number called weight, and a single bias. Applying a filter over a portion of the input volume consist in multiplying each element of the input with the corresponding weight, then all the products are added together, along with the bias term, and the sum is used as input to the non-linear activation function producing the final single-valued result. The filter is then slid along the dimenions of the input and the operations are repeated to form one channel of the feature map; stacking together the results of convolutions of multiple filters the complete feature map is obtained. Considering as an example the case of creating a convolutional layer for a RGB image, the input (the image) can be seen as a 2 dimensional volume with 3 channels. The filter will therefore be 3 dimensional and slide along the 2 dimensions of the input, and the result of convoluting the filter will be a 2 dimensional output. The size of the first 2 dimensions of the filter, who both affect how big the portion of the input captured by one applicaton is, can be chosen freely, but the third dimension needs to be less than or equal to the number of channels of the input, otherwise applying the filter would not be possible. In order for the sizes of the output dimensions to be equal to the input ones the input can be padded with zero valued elements. Adding more equally dimensioned filters to the layer increases the number of channels the output volume will have.

The parameters of a convolutional layer are therefore given by the weights and biases of each filter. During training the network can learn which parameters allow filters to extract the most relevant features from the input.

The pooling layer function is to reduce the spatial size of a feature map, practically downsampling the convolutional layer output in order to decrease the total number of parameters and the computational burden of the network. It does so by sliding a particular kind of filter, that returns the maximum value among the ones it's being applied on, across the dimensions of each input channel independently. The output produced has the same number of channels as the input but the size of the other dimensions is reduced, this way only the most dominant features detected by the convolutional layer are retained. Going back to the RGB image example, assuming that a convolutional layer used 10 filters and produced a feature map where the first and second dimensions are equal to 30, a pooling layer, whose filter is of size 2 on both dimension and slides across with a stride of 2 elements, outputs a volume of size 15 on both dimension and 10 channels.

A typical ConvNet combines a variable number of blocks composed of one or more convolutional layers followed by a pooling layer (it is to note that if the sizes of the feature map produced by the convolutional layer are already small, the pooling layer may be avoided). The last output volume of this chain is then flattened into a one dimensional vector, losing its grid-like structure, and connected to a feedforward network that handles the classification part. This allows to have the same loss function as a feedforward network and to similarly train the weights via stochastic gradient descent. Figure 2 shows one of the possible structures of a ConvNet.

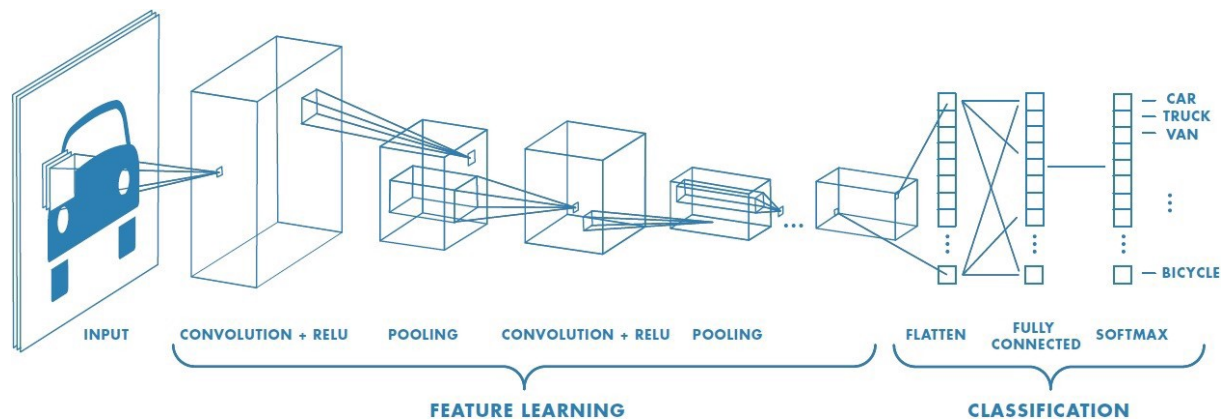


Figure 2: A visual representation of a convolutional neural network used for image classification. The main components are: the convolutional layers producing the feature maps, the pooling layers used to downsample the convolutional layers output, and a standard fully connected feedforward network to handle classification

Regularization consists in setting constraints to the model in order to avoid overfitting of the training set. Several techniques have been developed: L2 regularization adds to the loss function a fraction of the L2 norm of the weights, penalizing models where some of the weights have a value much greater than the rest. This is done to increase the generalization capabilities by encouraging the model to use all the input values when making a prediction, instead of having few input dimensions heavily influencing the prediction. Dropout consists in randomly (with a fixed probability) dropping neurons, along with their connections, from the network during training, preventing neurons from co-adapting too much. At test time all neurons are active and their outputs are scaled to match the output during training in expectation.

Lastly batch normalization is seen as a normalization technique that normalizes (mean subtraction and standard deviation division) the output of each activation function with respect to the values produced by that specific activation function for the batch that the network is iterating over during training, the normalized value is then linearly transformed (multiplication by a scalar and sum to a bias) to give the output a new distribution. The parameters, the scalar and the bias, that determine this transformation are learnt through training. Batch normalization has been shown to improve the prediction performance, even though the interpretation of its operations is not yet completely clear.

The dataset

The Street View House Number dataset consists of real world RGB images obtained from house numbers in Google Street View. The images are available in two formats: the original, variable-resolution, images, with a separate file describing the position of each digit in the image (shown in Figure 3a), and 32 by 32 pixels images centered around a single digit, an example of the second format is shown in Figure 3b.

The second format, with the images already cropped and centered around the digit to predict, is certainly to prefer for a classification problem. This is due to the fact that the input are all of the same shape and the format is similar to other datasets that have been extensively analyzed. It is to note that the preprocessing done leaves nonetheless other digits in the image that are nonrelevant to the prediction, from this a greater degree of difficulty than just predicting a single isolated digit arises. The faced problem is therefore a classification problem where the input are 32 by 32 pixels images classified into 10 classes, where an image representing a 9 has a numerical label 9, 1 has a numerical label 1 and so on, except for 0 which has label 10 in the dataset. The training set contains 73257 images and the test set 26032. Additional 531131, somewhat less difficult, images are available to use as extra training images, but they will not be



(a) The first format



(b) The second format

Figure 3: (a) The original images obtained from Street View. The blue rectangles describe the position of each digit in the images, they are shown for clarity but are not present in the images, their coordinates are stored in a separate file. (b) The images cropped and centered around a single digit. Note the presence of other digits and distracting elements to the sides of the one centered.

used in this experiment. A validation set for hyperparameter tuning has been randomly extracted from the training set, the validation set size is 10% of the original training set size; this resulted in a smaller training set during the model design phase. The distribution of the labels for the data sets is shown in Figure 4. Because of the nature of house numbering it is expected to be present a higher count of lower digits, particularly 1 and 2.

The data set can be obtained at the following website: <http://ufldl.stanford.edu/housenumbers/>

Developing the model

The first step involved some data preprocessing. The images are colored but the colors do not seem to convey any useful information that can be used to distinguish one digit from another, therefore the images have been converted from RGB to grayscale using the formula $gv =$




Figure 4: Histograms representing the absolute frequency of the labels for the training set (a), validation set (b), and test set (c). The distribution is very similar among the various sets, making it sound to use them for machine learning tasks.

$0.2989r + 0.5870g + 0.1140b$ where r, g, b are the pixel values for the red, green, and blue channels and gv is the corresponding grayscale value of that pixel. This was followed by a mean subtraction and a normalization. Mean subtraction consists in subtracting the mean across every feature of the input, making them zero centered; for images, in particular, the mean among all the pixel in the training set can be calculated and subtracted instead. The normalization is obtained by dividing each pixel by the standard deviation of all the pixel after they have been zero centered. It is important to notice that the mean and the standard deviation have to be calculated on the training set only and then used to preprocess the training set, the validation set, and the test set. This is due to the fact that, in a real world scenario, the model can't have those statistics for future, never seen before, data points it will have to classify, and since the test set is supposed to approximate those data points its statistics cannot be used.

In order to design the model and tune hyperparameters it has been used a form of cross validation consisting of the following steps: select a starting value for each hyperparameter, choose a hyperparameter to optimize, perform a grid search on that hyperparameter using the loss on the validation set as the search metric (training is done on the reduced training set consisting of the original training set from which examples added to the validation set have been removed), consider this hyperparameter optimized and repeat these steps on the remaining not already optimized hyperparameters. It is worth noting that this does not represent a proper hyperparameter search because the order in which the hyperparameters are optimized affects the optimization. A more correct way to perform the search, albeit much more computationally expensive, would be to train the model with numerous combinations of values of the hyperparameters and choosing the combination of values that performs best on the validation set. This method has not been performed in this experiment due to its computational burden.

As described in the first section of this report a ConvNet has typically three components:

convolutional layers, pooling layers and a feedforward component. The base network architecture that has been proven to have good results for similar tasks consists of blocks made of one or more convolutional layers followed by a pooling layer, these blocks are concatenated until the output of the last pooling layer has a small enough size. At this point one or more convolutional layers can follow to increase the representational power of the network. The resulting output is then flattened into a vector and fed to a feedforward network with a softmax output layer for classification purpose. The filters of the convolutional layers have size equal to 3 on each dimension, these small filters have been shown to consistently perform better than larger filters, especially when stacking multiple convolutional layers on top of each other. The input of convolutional layers is zero-padded before performing the convolution so that the sizes of the output dimensions remain unchanged. The filters of the pooling layers have size equal to 2 for each dimension, halving the size of the input dimensions; having larger filters would discard too much information. Considering the original input size is 32 by 32, no more than two pooling layers can be applied, otherwise the output would be smaller than 8 by 8 which would be too small to convey enough information for the classification.

In the first section it was stated that parameters of the network can be updated using stochastic gradient descent which follows the update rule (4) where the learning rate $\eta^{(t)}$ at timestep t is a hyperparameter. Adam, short for Adaptive Moment Estimation, is used in this experiment instead of gradient descent. It is a similar but more recent technique that improves convergence time and adjusts the learning rate for each parameter based on the update rule:

$$\omega^{(t+1)} \leftarrow \omega^{(t)} - \eta^{(t)} \frac{\hat{\mathbf{m}}_\omega}{\sqrt{\hat{\mathbf{v}}_\omega} + \epsilon}$$

where $\omega^{(t)}$ is the vector of parameters at time step t , ϵ is a small constant to avoid division by zero, and $\hat{\mathbf{m}}_\omega$ and $\hat{\mathbf{v}}_\omega$ are, respectively, estimates of the first and second moment of the gradient, and are given by:

$$\begin{aligned} \hat{\mathbf{m}}_\omega &= \frac{\mathbf{m}_\omega^{(t)}}{1 - \beta_1^t}; & \mathbf{m}_\omega^{(t)} &\leftarrow \beta_1 \mathbf{m}_\omega^{(t-1)} + (1 - \beta_1) \nabla_\omega \ell^{(t)} \\ \hat{\mathbf{v}}_\omega &= \frac{\mathbf{v}_\omega^{(t)}}{1 - \beta_2^t}; & \mathbf{v}_\omega^{(t)} &\leftarrow \beta_2 \mathbf{v}_\omega^{(t-1)} + (1 - \beta_2) (\nabla_\omega \ell^{(t)})^2 \end{aligned}$$

β_1 and β_2 are hyperparameters called forgetting factors and usually, as well as in this experiment, they assume values of 0.9 and 0.999 respectively, $\mathbf{m}_\omega^{(t)}$ and $\mathbf{v}_\omega^{(t)}$ are equal to the zero vector for $t = 0$; squaring and square-rooting is done elementwise. This means that the learning rate $\eta^{(t)}$ is divided by a factor $\hat{\mathbf{v}}_\omega$ proportional to the moving average of the square of the gradient

hyperparameter	starting	validated	enlarged
conv layers 1st block	1	3	7
conv layers 2nd block	1	3	6
conv layers 3rd block	1	3	5
filters 1st block	32	64	128
filters 2nd block	64	128	256
filters 3rd block	64	128	384
feedforward layer	1	2	2
neurons 1st layer	128	96	352
neurons 2nd layer	off	64	192
batch size	128	96	96
learning rate	7E-4	6.5E-4	6.5E-4
learning rate decay	off	on	on
dropout	off	0.2	0.2
batch normalization	off	on	on
l2 regularization	off	off	off

Table 1: hyperparameters and their values at each phase of the training

over time, the bigger the average the smaller the learning rate, and the step \hat{m}_ω is proportional to the running average of the gradient over time, therefore the last gradient influences only a fraction of the update.

The starting hyperparameters and their respective values are given in Table 1. These hyperparameter were validate, as describe above, in this order (a block corresponds to a series of convolutional layers followed by a pooling layer): number of convolutional layers in each block, number of filters in each layer of each block, number of layers in the feedforward component, number of neurons in the feedforward layers, applying dropout in the feedforward layers, batch size, l2 regularization, learning rate, learning rate decay, and batch normalization. After the validation the model was enlarged by adding more parameters to it as long as the loss function on the validation set decreased. The validated values of the hyperparameters, as well as the values of the enlarged model, are shown in Table 1. The learning rate decays exponentially following the rule: $\eta = (10^{-3})(0.65)^p$, where p is the epoch number.

The value of the loss function on the validation set decreased throughout the steps going from the value of 0.37 at the start, to 0.21 after validation, to 0.19 after enlargement. Likewise accuracy increased from 93.2% and 89.2% to 99.9% and 95.4% on training and validation set respectively. It is worth noting that the model converges extremely quickly to the best solution after about five epochs, then slight overfitting arises and increasing regularization factors don't improve this aspect. Graphs for accuracy and loss during training of the enlarged model are

shown in Figure 5. In order to achieve better results a more complex network architecture would probably be necessary.




Figure 5: Loss and accuracy on the validation and reduced training set of the final enlarged model after each epoch of training on the reduced training set.

Results

The final model (called enlarged in Table 1) has been trained on the entire training set and evaluated on the test set. The model performed substantially better during testing obtaining a loss of 0.15 and accuracy of 96.2% on the test set. This increase of performance may be attributed to the larger training set compared to the validation phase. Graphs of the loss and accuracy of the model during training are given in Figure 6.




Figure 6: Loss and accuracy on the test and training set of the final enlarged model after each epoch of training on the complete training set

The confusion matrix of the test set predictions, provided in Figure 7, shows that the model tends to misclassify the digit 7 as 1 more often compared to the other errors. This may be

interpreted in a positive way meaning that, since the number 7 and 1 have relatively similar shapes, the model is actually learning the shapes of the digits. The same argument can be made for other frequent misclassification such as the digit 8 predicted as 6 and the digit 6 predicted as 5.




Figure 7: The confusion matrix given by the predictions of the model on the test set. The row denotes the true label, the column the predicted label, and the value is the percentages of data points with true label given by the row and prediction given by the column.

The model uses the *softmax* function in the last output layers, because of this its outputs can be interpreted as a distribution probability over the labels given a data point (as explained in the first section) and the prediction will be the label with the highest probability. This allows to seek among the wrong predictions the ones that the model classified with the highest confidence (highest probability value). These predictions revealed that there are instances of mislabeled examples in the test set. Figure 8 shows the forty wrong predictions with the highest confidence and clearly the majority of those data points are mislabeled. A similar analysis among the wrong predictions with the lowest confidence didn't show any kind of mislabeling of those examples. It remains unclear the full extent to which the mislabeling affects the test set and possibly the training set, and consequently the evaluation and training of the model.



Figure 8: The top forty wrong predictions to which the model assigns the highest confidence. It is clear the presence of some mislabeling of the examples. The images are shown after the conversion from RGB to grayscale.

In conclusion the model classified the test set with an accuracy of 96.2% and reached convergence extremely quickly in about 7 epochs. Possible ways to improve these results would be to apply typical image processing methods, such as contrast normalization and luminosity normalization, to preprocess the data points, to use a more complex neural network architecture, to understand to which extent, if at all, the mislabeling affects the training of the model, to increase the training set size by using the trivial examples, and, finally, to perform a more thorough hyperparameter search.